

CNN ARCHITECTURES FOR LARGE-SCALE AUDIO CLASSIFICATION

Shawn Hershey, Sourish Chaudhuri, Daniel P. W. Ellis, Jort F. Gemmeke, Aren Jansen, R. Channing Moore,
Manoj Plakal, Devin Platt, Rif A. Saurous, Bryan Seybold, Malcolm Slaney, Ron J. Weiss, Kevin Wilson

Google, Inc., New York, NY, and Mountain View, CA, USA

shershey@google.com

ABSTRACT

Convolutional Neural Networks (CNNs) have proven very effective in image classification and show promise for audio. We use various CNN architectures to classify the soundtracks of a dataset of 70M training videos (5.24 million hours) with 30,871 video-level labels. We examine fully connected Deep Neural Networks (DNNs), AlexNet [1], VGG [2], Inception [3], and ResNet [4]. We investigate varying the size of both training set and label vocabulary, finding that analogs of the CNNs used in image classification do well on our audio classification task, and larger training and label sets help up to a point. A model using embeddings from these classifiers does much better than raw features on the *Audio Set* [5] Acoustic Event Detection (AED) classification task.

Index Terms— Acoustic Event Detection, Acoustic Scene Classification, Convolutional Neural Networks, Deep Neural Networks, Video Classification

1. INTRODUCTION

Image classification performance has improved greatly with the advent of large datasets such as ImageNet [6] using Convolutional Neural Network (CNN) architectures such as AlexNet [1], VGG [2], Inception [3], and ResNet [4]. We are curious to see if similarly large datasets and CNNs can yield good performance on audio classification problems. Our dataset consists of 70 million (henceforth 70M) training videos totalling 5.24 million hours, each tagged from a set of 30,871 (henceforth 30K) labels. We call this dataset YouTube-100M. Our primary task is to predict the video-level labels using audio information (i.e., soundtrack classification). Per Lyon [7], teaching machines to hear and understand video can improve our ability to “categorize, organize, and index them”.

In this paper, we use the YouTube-100M dataset to investigate: how popular Deep Neural Network (DNN) architectures compare on video soundtrack classification; how performance varies with different training set and label vocabulary sizes; and whether our trained models can also be useful for AED.

Historically, AED has been addressed with features such as MFCCs and classifiers based on GMMs, HMMs, NMF, or SVMs [8, 9, 10, 11]. More recent approaches use some form of DNN, including CNNs [12] and RNNs [13]. Prior work has been reported on datasets such as TRECVID [14], ActivityNet [15], Sports1M [16], and TUT/DCASE Acoustic scenes 2016 [17] which are much smaller than YouTube-100M. Our large dataset puts us in a good position to evaluate models with large model capacity.

RNNs and CNNs have been used in Large Vocabulary Continuous Speech Recognition (LVCSR) [18]. Unlike that task, our labels apply to entire videos without any changes in time, so we have yet to try such recurrent models.

Eghbal-Zadeh et al. [19] recently won the DCASE 2016 Acoustic Scene Classification (ASC) task, which, like soundtrack classification, involves assigning a single label to an audio clip containing many events. Their system used spectrogram features feeding a VGG classifier, similar to one of the classifiers in our work. This paper, however, compares the performance of several different architectures. To our knowledge, we are the first to publish results of Inception and ResNet networks applied to audio.

We aggregate local classifications to whole-soundtrack decisions by imitating the visual-based video classification of Ng et al. [20]. After investigating several more complex models for combining information across time, they found simple averaging of single-frame CNN classification outputs performed nearly as well. By analogy, we apply a classifier to a series of non-overlapping segments, then average all the sets of classifier outputs.

Kumar et al. [21] consider AED in a dataset with video-level labels as a Multiple Instance Learning (MIL) problem, but remark that scaling such approaches remains an open problem. By contrast, we are investigating the limits of training with weak labels for very large datasets. While many of the individual segments will be uninformative about the labels inherited from the parent video, we hope that, given enough training, the net can learn to spot useful cues. We are not able to quantify how “weak” the labels are (i.e., what proportion of the segments are uninformative), and for the majority of classes (e.g., “Computer Hardware”, “Boeing 757”, “Ollie”), it’s not clear how to decide relevance. Note that for some classes (e.g. “Beach”), background ambiance is itself informative.

Our dataset size allows us to examine networks with large model capacity, fully exploiting ideas from the image classification literature. By computing log-mel spectrograms of multiple frames, we create 2D image-like patches to present to the classifiers. Although the distinct meanings of time and frequency axes might argue for audio-specific architectures, this work employs minimally-altered image classification networks such as Inception-V3 and ResNet-50. We train with subsets of YouTube-100M spanning 23K to 70M videos to evaluate the impact of training set size on performance, and we investigate the effects of label set size on generalization by training models with subsets of labels, spanning 400 to 30K, which are then evaluated on a single common subset of labels. We additionally examine the usefulness of our networks for AED by examining the performance of a model trained with embeddings from one of our networks on the *Audio Set* [5] dataset.

2. DATASET

The YouTube-100M data set consists of 100 million YouTube videos: 70M training videos, 10M evaluation videos, and a pool of 20M videos that we use for validation. Videos average 4.6 minutes each for a total of 5.4M training hours. Each of these videos

Table 1: Example labels from the 30K set.

| Label prior | Example Labels |
|----------------|----------------------------------------|
| 0.1 . . . 0.2 | Song, Music, Game, Sports, Performance |
| 0.01 . . . 0.1 | Singing, Car, Chordophone, Speech |
| $\sim 10^{-5}$ | Custom Motorcycle, Retaining Wall |
| $\sim 10^{-6}$ | Cormorant, Lecturer |

is labeled with 1 or more topic identifiers (from Knowledge Graph [22]) from a set of 30,871 labels. There are an average of around 5 labels per video. The labels are assigned automatically based on a combination of metadata (title, description, comments, etc.), context, and image content for each video. The labels apply to the entire video and range from very generic (e.g. “Song”) to very specific (e.g. “Cormorant”). Table 1 shows a few examples.

Being machine generated, the labels are not 100% accurate and of the 30K labels, some are clearly acoustically relevant (“Trumpet”) and others are less so (“Web Page”). Videos often bear annotations with multiple degrees of specificity. For example, videos labeled with “Trumpet” are often labeled “Entertainment” as well, although no hierarchy is enforced.

3. EXPERIMENTAL FRAMEWORK

3.1. Training

The audio is divided into non-overlapping 960 ms frames. This gave approximately 20 billion examples from the 70M videos. Each frame inherits all the labels of its parent video. The 960 ms frames are decomposed with a short-time Fourier transform applying 25 ms windows every 10 ms. The resulting spectrogram is integrated into 64 mel-spaced frequency bins, and the magnitude of each bin is log-transformed after adding a small offset to avoid numerical issues. This gives log-mel spectrogram patches of 96×64 bins that form the input to all classifiers. During training we fetch mini-batches of 128 input examples by randomly sampling from all patches.

All experiments used TensorFlow [23] and were trained asynchronously on multiple GPUs using the Adam [24] optimizer. We performed grid searches over learning rates, batch sizes, number of GPUs, and parameter servers. Batch normalization [25] was applied after all convolutional layers. All models used a final sigmoid layer rather than a softmax layer since each example can have multiple labels. Cross-entropy was the loss function. In view of the large training set size, we did not use dropout [26], weight decay, or other common regularization techniques. For the models trained on 7M or more examples, we saw no evidence of overfitting. During training, we monitored progress via 1-best accuracy and mean Average Precision (mAP) over a validation subset.

3.2. Evaluation

From the pool of 10M evaluation videos we created three balanced evaluation sets, each with roughly 33 examples per class: 1M videos for the 30K labels, 100K videos for the 3087 (henceforth 3K) most frequent labels, and 12K videos for the 400 most frequent labels. We passed each 960 ms frame from each evaluation video through the classifier. We then averaged the classifier output scores across all segments in a video.

For our metrics, we calculated the balanced average across all classes of AUC (also reported as the equivalent d-prime class separation), and mean Average Precision (mAP). AUC is the area under the Receiver Operating Characteristic (ROC) curve [27], that is, the

probability of correctly classifying a positive example (correct accept rate) as a function of the probability of incorrectly classifying a negative example as positive (false accept rate); perfect classification achieves AUC of 1.0 (corresponding to an infinite d-prime), and random guessing gives an AUC of 0.5 (d-prime of zero).¹ mAP is the mean across classes of the Average Precision (AP), which is the proportion of positive items in a ranked list of trials (i.e., Precision) averaged across lists just long enough to include each individual positive trial [28]. AP is widely used as an indicator of precision that does not require a particular retrieval list length, but, unlike AUC, it is directly correlated with the prior probability of the class. Because most of our classes have very low priors ($< 10^{-4}$), the mAPs we report are typically small, even though the false alarm rates are good.

3.3. Architectures

Our baseline is a fully connected DNN, which we compared to several networks closely modeled on successful image classifiers. For our baseline experiments, we trained and evaluated using only the 10% most frequent labels of the original 30K (i.e., 3K labels). For each experiment, we coarsely optimized number of GPUs and learning rate for the frame level classification accuracy. The optimal number of GPUs represents a compromise between overall computing power and communication overhead, and varies by architecture.

3.3.1. Fully Connected

Our baseline network is a fully connected model with RELU activations [29], N layers, and M units per layer. We swept over $N = [2, 3, 4, 5, 6]$ and $M = [500, 1000, 2000, 3000, 4000]$. Our best performing model had $N = 3$ layers, $M = 1000$ units, learning rate of 3×10^{-5} , 10 GPUs and 5 parameter servers. This network has approximately 11.2M weights and 11.2M multiplies.

3.3.2. AlexNet

The original AlexNet [1] architectures was designed for a $224 \times 224 \times 3$ input with an initial 11×11 convolutional layer with a stride of 4. Because our inputs are 96×64 , we use a stride of 2×1 so that the number of activations are similar after the initial layer. We also use batch normalization after each convolutional layer instead of local response normalization (LRN) and replace the final 1000 unit layer with a 3087 unit layer. While the original AlexNet has approximately 62.4M weights and 1.1G multiplies, our version has 37.3M weights and 767M multiplies. Also, for simplicity, unlike the original AlexNet, we do not split filters across multiple devices. We trained with 20 GPUs and 10 parameter servers.

3.3.3. VGG

The only changes we made to VGG (configuration E) [2] were to the final layer (3087 units with a sigmoid) as well as the use of batch normalization instead of LRN. While the original network had 144M weights and 20B multiplies, the audio variant uses 62M weights and 2.4B multiplies. We tried another variant that reduced the initial strides (as we did with AlexNet), but found that not modifying the strides resulted in faster training and better performance. With our setup, parallelizing beyond 10 GPUs did not help significantly, so we trained with 10 GPUs and 5 parameter servers.

¹ $d' = \sqrt{2}F^{-1}(\text{AUC})$ where F^{-1} is the inverse cumulative distribution function for a unit Gaussian.

Table 2: Comparison of performance of several DNN architectures trained on 70M videos, each tagged with labels from a set of 3K. The last row contains results for a model that was trained much longer than the others, with a reduction in learning rate after 13 million steps.

| Architectures | Steps | Time | AUC | d-prime | mAP |
|-----------------|-------|------|--------------|--------------|--------------|
| Fully Connected | 5M | 35h | 0.851 | 1.471 | 0.058 |
| AlexNet | 5M | 82h | 0.894 | 1.764 | 0.115 |
| VGG | 5M | 184h | 0.911 | 1.909 | 0.161 |
| Inception V3 | 5M | 137h | 0.918 | 1.969 | 0.181 |
| ResNet-50 | 5M | 119h | 0.916 | 1.952 | 0.182 |
| ResNet-50 | 17M | 356h | 0.926 | 2.041 | 0.212 |

3.3.4. Inception V3

We modified the inception V3 [3] network by removing the first four layers of the stem, up to and including the MaxPool, as well as removing the auxiliary network. We changed the Average Pool size to 10×6 to reflect the change in activations. We tried including the stem and removing the first stride of 2 and MaxPool but found that it performed worse than the variant with the truncated stem. The original network has 27M weights with 5.6B multiplies, and the audio variant has 28M weights and 4.7B multiplies. We trained with 40 GPUs and 20 parameter servers.

3.3.5. ResNet-50

We modified ResNet-50 [4] by removing the stride of 2 from the first 7×7 convolution so that the number of activations was not too different in the audio version. We changed the Average Pool size to 6×4 to reflect the change in activations. The original network has 26M weights and 3.8B multiplies. The audio variant has 30M weights and 1.9B multiplies. We trained with 20 GPUs and 10 parameter servers.

4. EXPERIMENTS

4.1. Architecture Comparison

For all network architectures we trained with 3K labels and 70M videos and compared after 5 million mini-batches of 128 inputs. Because some networks trained faster than others, comparing after a fixed wall-clock time would give slightly different results but would not change the relative ordering of the architectures' performance. We include numbers for ResNet after training for 17 million mini-batches (405 hours) to show that performance continues to improve. We reduced the learning rate by a factor of 10 after 13 million mini-batches.

Table 2 shows the evaluation results calculated over the 100K balanced videos. All CNNs beat the fully-connected baseline. Inception and ResNet achieve the best performance; they provide high model capacity and their convolutional units can efficiently capture common structures that may occur in different areas of the input array for both images, and, we infer, our audio representation.

To investigate how the prior likelihood of each label affects its performance, Fig. 1 shows a scatter plot of the 30K classes with label frequency on the x axis and ResNet-50's d-prime on the y axis. d-prime seems to stay centered around 2.0 across label prior, although the variance of d-prime increases for less-common classes. This is contrary to the usual result where classifier performance improves with increased training data, particularly over the 5 orders of magnitude illustrated in the plot.




Fig. 1: Scatter plot of ResNet-50's per-class d-prime versus log prior probability. Each point is a separate class from a random 20% subset of the 30K set. Color reflects the class AP.

4.2. Label Set Size

Using a 400 label subset of the 30K labels, we investigated how training with different subsets of classes can affect performance, perhaps by encouraging intermediate representations that better generalize to unseen examples even for the evaluation classes. In addition to examining three label set sizes (30K, 3K, and 400), we compared models with and without a bottleneck layer of 128 units placed right before the final output layer. We introduced the bottleneck layer to speed up the training of the model trained with 30K labels. Without a bottleneck, the larger output layer increased the number of weights from 30M to 80M and significantly reduced training speed. We do not report metrics on the 30K label model without the bottleneck because it would have taken several months to train. For all label set size experiments, we used the ResNet-50 model and trained for 5 million mini-batches of 128 inputs (about 120 hours) on 70M videos.

Tables 3 shows the results. When comparing models with the bottleneck, we see that performance does indeed improve slightly as we increase the number of labels we trained on, although networks without the bottleneck have higher performance overall. The bottleneck layer is relatively small compared to the 2048 activations coming out of ResNet-50's Average Pool layer and so it is effecting a substantial reduction in information. These results provide weak support to the notion that training with a broader set of categories can help to regularize even the 400 class subset.

4.3. Training Set Size

Having a very large training set available allows us to investigate how training set size affects performance. With 70M videos and an average of 4.6 minutes per video, we have around 20 billion 960 ms training examples. Given ResNet-50's training speed of 11 mini-batches per second with 20 GPUs, it would take 23 weeks for the network to see each pattern once (one epoch). However, if all videos were equal length and fully randomized, we expect to see at least one frame from each video in only 14 hours. We hypothesize that, even

Table 3: Results of varying label set size, evaluated over 400 labels. All models are variants of ResNet-50 trained on 70M videos. The bottleneck, if present, is 128 dimensions.

| Bneck | Labels | AUC | d-prime | mAP |
|-------|--------|--------------|--------------|--------------|
| no | 30K | — | — | — |
| no | 3K | 0.930 | 2.087 | 0.381 |
| no | 400 | 0.928 | 2.067 | 0.376 |
| yes | 30K | 0.925 | 2.035 | 0.369 |
| yes | 3K | 0.919 | 1.982 | 0.347 |
| yes | 400 | 0.924 | 2.026 | 0.365 |




Fig. 2: Three example excerpts from a video classified by ResNet-50 with instantaneous model outputs overlaid. The 16 classifier outputs with the greatest peak values across the entire video were chosen from the 30K set for display.

if we cannot get through an entire epoch, 70M videos will provide an advantage over 7M by virtue of the greater diversity of videos underlying the limited number of training patterns consumed. We trained a ResNet-50 model for 16 million mini-batches of 128 inputs (about 380 hours) on the 3K label set with 70M, 7M, 700K, 70K, and 23K videos.

The video level results are shown in Table 4. The 70K and 23K models show worse performance but the validation plots (not included) showed that they likely suffered from overfitting. Regularization techniques (or data augmentation) might have boosted the numbers on these smaller training sets. The 700K, 7M, and 70M models are mostly very close in performance although the 700K model is slightly inferior.

4.4. AED with the Audio Set Dataset

Audio Set [5] is a dataset of over 1 million 10 second excerpts labeled with a vocabulary of acoustic events (whereas not all of the YouTube-100M 30K labels pertain to acoustic events). This comes to about 3000 hours – still only $\approx 0.05\%$ of YouTube-100M. We train two fully-connected models to predict labels for *Audio Set*. The first model uses 64×20 log-mel patches and the second uses the output of the penultimate “embedding” layer of our best ResNet model as inputs. The log-mel baseline achieves a balanced mAP of 0.137 and AUC of 0.904 (equivalent to d-prime of 1.846). The model trained on embeddings achieves mAP / AUC / d-prime of 0.314 / 0.959 / 2.452. This jump in performance reflects the benefit of the larger YouTube-100M training set embodied in the ResNet classifier outputs.

Table 4: Results of training with different amounts of data. All rows used the same ResNet-50 architecture trained on videos tagged with labels from a set of 3K.

| Training Videos | AUC | d-prime | mAP |
|-----------------|--------------|--------------|--------------|
| 70M | 0.923 | 2.019 | 0.206 |
| 7M | 0.922 | 2.006 | 0.202 |
| 700K | 0.921 | 1.997 | 0.203 |
| 70K | 0.909 | 1.883 | 0.162 |
| 23K | 0.868 | 1.581 | 0.118 |

5. CONCLUSIONS

The results in Section 4.1 show that state-of-the-art image networks are capable of excellent results on audio classification when compared to a simple fully connected network or earlier image classification architectures. In Section 4.2 we saw results showing that training on larger label set vocabularies can improve performance, albeit modestly, when evaluating on smaller label sets. In Section 4.3 we saw that increasing the number of videos up to 7M improves performance for the best-performing ResNet-50 architecture. We note that regularization could have reduced the gap between the models trained on smaller datasets and the 7M and 70M datasets. In Section 4.4 we see a significant increase over our baseline when training a model for AED with ResNet embeddings on the *Audio Set* dataset.

In addition to these quantified results, we can subjectively examine the performance of the model on segments of video. Fig. 2 shows the results of running our best classifier over a video and overlaying the frame-by-frame results of the 16 classifier outputs with the greatest peak values across the entire video. The different sound sources present at different points in the video are clearly distinguished.²

6. ACKNOWLEDGEMENTS

The authors would like to thank George Toderici and Marvin Ritter, both with Google, for their very valuable feedback.

²A similar video is available online at https://youtu.be/oAAo_r7ZT8U.

7. REFERENCES

- [1] A. Krizhevsky, I. Sutskever, and G. E. Hinton, “Imagenet classification with deep convolutional neural networks,” in *Advances in neural information processing systems*, 2012, pp. 1097–1105.
- [2] K. Simonyan and A. Zisserman, “Very deep convolutional networks for large-scale image recognition,” *arXiv preprint arXiv:1409.1556*, 2014.
- [3] C. Szegedy, V. Vanhoucke, S. Ioffe, J. Shlens, and Z. Wojna, “Rethinking the inception architecture for computer vision,” *arXiv preprint arXiv:1512.00567*, 2015.
- [4] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” *arXiv preprint arXiv:1512.03385*, 2015.
- [5] J. F. Gemmeke, D. P. W. Ellis, D. Freedman, A. Jansen, W. Lawrence, R. C. Moore, M. Plakal, and M. Ritter, “Audio Set: An ontology and human-labeled dataset for audio events,” in *IEEE ICASSP 2017*, New Orleans, 2017.
- [6] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei, “Imagenet: A large-scale hierarchical image database,” in *Computer Vision and Pattern Recognition, 2009. CVPR 2009. IEEE Conference on*. IEEE, 2009, pp. 248–255.
- [7] R. F. Lyon, “Machine hearing: An emerging field [exploratory dsp],” *Ieee signal processing magazine*, vol. 27, no. 5, pp. 131–139, 2010.
- [8] A. Mesaros, T. Heittola, A. Eronen, and T. Virtanen, “Acoustic event detection in real life recordings,” in *Signal Processing Conference, 2010 18th European*. IEEE, 2010, pp. 1267–1271.
- [9] X. Zhuang, X. Zhou, M. A. Hasegawa-Johnson, and T. S. Huang, “Real-world acoustic event detection,” *Pattern Recognition Letters*, vol. 31, no. 12, pp. 1543–1551, 2010.
- [10] J. F. Gemmeke, L. Vugtgen, P. Karsmakers, B. Vanrumste, et al., “An exemplar-based nmf approach to audio event detection,” in *2013 IEEE Workshop on Applications of Signal Processing to Audio and Acoustics*. IEEE, 2013, pp. 1–4.
- [11] A. Temko, R. Malkin, C. Zieger, D. Macho, C. Nadeu, and M. Omologo, “Clear evaluation of acoustic event detection and classification systems,” in *International Evaluation Workshop on Classification of Events, Activities and Relationships*. Springer, 2006, pp. 311–322.
- [12] N. Takahashi, M. Gygli, B. Pfister, and L. Van Gool, “Deep convolutional neural networks and data augmentation for acoustic event detection,” *arXiv preprint arXiv:1604.07160*, 2016.
- [13] G. Parascandolo, H. Huttunen, and T. Virtanen, “Recurrent neural networks for polyphonic sound event detection in real life recordings,” in *2016 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2016, pp. 6440–6444.
- [14] G. Awad, J. Fiscus, M. Michel, D. Joy, W. Kraaij, A. F. Smeaton, G. Quenot, M. Eskevich, R. Aly, and R. Ordelman, “Trecvid 2016: Evaluating video search, video event detection, localization, and hyperlinking,” in *Proceedings of TRECVID 2016*. NIST, USA, 2016.
- [15] B. G. Fabian Caba Heilbron, Victor Escorcia and J. C. Niebles, “Activitynet: A large-scale video benchmark for human activity understanding,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2015, pp. 961–970.
- [16] A. Karpathy, G. Toderici, S. Shetty, T. Leung, R. Sukthankar, and L. Fei-Fei, “Large-scale video classification with convolutional neural networks,” in *CVPR*, 2014.
- [17] A. Mesaros, T. Heittola, and T. Virtanen, “TUT database for acoustic scene classification and sound event detection,” in *24th European Signal Processing Conference 2016 (EUSIPCO 2016)*, Budapest, Hungary, 2016, <http://www.cs.tut.fi/sgn/arg/dcase2016/>.
- [18] T. N. Sainath, O. Vinyals, A. Senior, and H. Sak, “Convolutional, long short-term memory, fully connected deep neural networks,” in *2015 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2015, pp. 4580–4584.
- [19] H. Eghbal-Zadeh, B. Lehner, M. Dorfer, and G. Widmer, “Cpjku submissions for dcase-2016: A hybrid approach using binaural i-vectors and deep convolutional neural networks.” .
- [20] J. Yue-Hei Ng, M. Hausknecht, S. Vijayanarasimhan, O. Vinyals, R. Monga, and G. Toderici, “Beyond short snippets: Deep networks for video classification,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2015, pp. 4694–4702.
- [21] A. Kumar and B. Raj, “Audio event detection using weakly labeled data,” *arXiv preprint arXiv:1605.02401*, 2016.
- [22] A. Singhal, “Introducing the knowledge graph: things, not strings,” 2012, Official Google blog, <https://googleblog.blogspot.com/2012/05/introducing-knowledge-graph-things-not.html>.
- [23] M. Abadi et al., “TensorFlow: Large-scale machine learning on heterogeneous systems,” 2015, Software available from tensorflow.org.
- [24] D. Kingma and J. Ba, “Adam: A method for stochastic optimization,” *arXiv preprint arXiv:1412.6980*, 2014.
- [25] S. Ioffe and C. Szegedy, “Batch normalization: Accelerating deep network training by reducing internal covariate shift,” *arXiv preprint arXiv:1502.03167*, 2015.
- [26] N. Srivastava, G. E. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov, “Dropout: a simple way to prevent neural networks from overfitting,” *Journal of Machine Learning Research*, vol. 15, no. 1, pp. 1929–1958, 2014.
- [27] T. Fawcett, “Roc graphs: Notes and practical considerations for researchers,” *Machine learning*, vol. 31, no. 1, pp. 1–38, 2004.
- [28] C. Buckley and E. M. Voorhees, “Retrieval evaluation with incomplete information,” in *Proceedings of the 27th annual international ACM SIGIR conference on Research and development in information retrieval*. ACM, 2004, pp. 25–32.
- [29] V. Nair and G. E. Hinton, “Rectified linear units improve restricted boltzmann machines,” in *Proceedings of the 27th International Conference on Machine Learning (ICML-10)*, 2010, pp. 807–814.